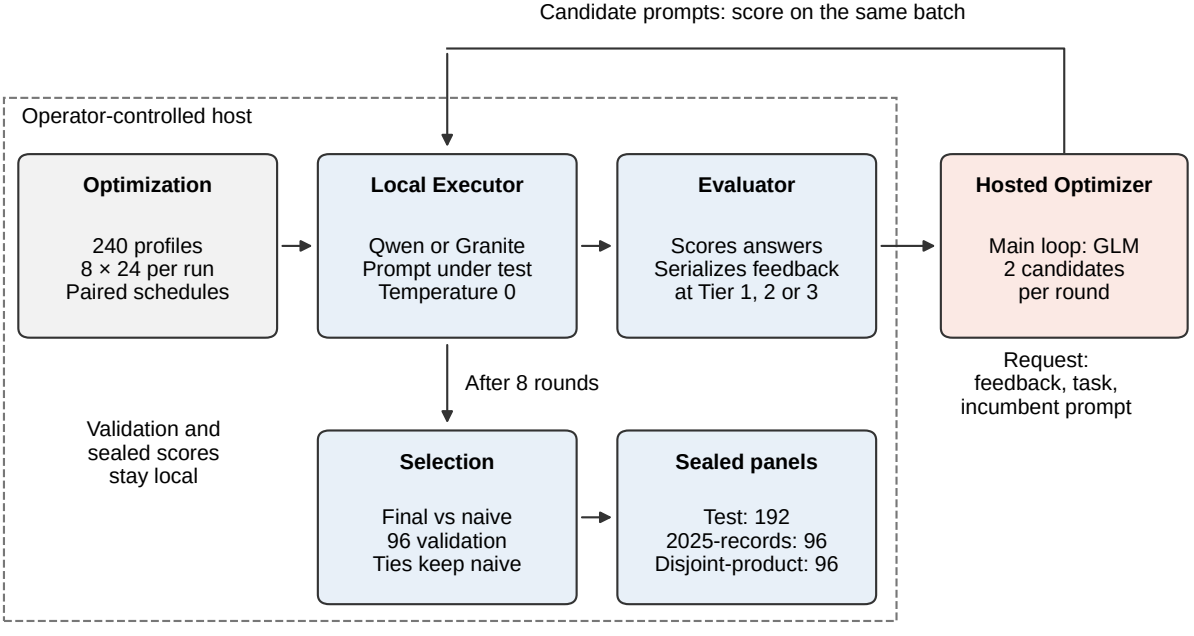
Candidate prompts: score on the same batch
Operator-controlled host
Optimization
240 profiles
8 × 24 per run
Paired schedules
Local Executor
Qwen or Granite
Prompt under test
Temperature 0
Evaluator
Scores answers
Serializes feedback
at Tier 1, 2 or 3
Hosted Optimizer
Main loop: GLM
2 candidates
per round
Request:
feedback, task,
incumbent prompt
After 8 rounds
Validation and
sealed scores
stay local
Selection
Final vs naive
96 validation
Ties keep naive
Sealed panels
Test: 192
2025-records: 96
Disjoint-product: 96