

Candidate prompts: score on the same batch

Operator-controlled host

Optimization

240 profiles
8 × 24 per run
Paired schedules

Local executor

Qwen or Granite
Prompt under test
Temperature 0

Evaluator

Scores answers
Serializes feedback
at Tier 1, 2 or 3

Hosted optimizer

Main loop: GLM
2 candidates
per round

Request:
feedback, task,
current prompt

After 8 rounds

Selection

Final vs naive
96 validation
Ties keep naive

Sealed panels

Test: 192
Temporal: 96
Held-out: 96

Validation and
sealed scores
stay local

